

Complexity of Hybrid Logics over Transitive Frames

Martin Mundhenk¹, Thomas Schneider¹, Thomas Schwentick², and
Volker Weber²

¹ Institut für Informatik, Friedrich-Schiller-Universität Jena, Germany
{mundhenk, schneider}@cs.uni-jena.de

² Fachbereich Informatik, Universität Dortmund, Germany
{thomas.schwentick, volker.weber}@udo.edu

Abstract. This paper examines the complexity of hybrid logics over transitive frames and transitive trees. We show that satisfiability over transitive frames for the hybrid language extended with \downarrow is NEXP-complete. This is in contrast to undecidability of satisfiability over arbitrary frames for this language [2]. We also show that adding @ or the past modality leads to undecidability over transitive frames, but not over transitive trees, where we show the richest language to be nonelementarily decidable. Moreover, we establish 2EXP and EXP upper bounds for satisfiability over transitive frames and transitive trees, respectively, for the hybrid *Until/Since* language. An EXP lower bound is shown to hold for the modal *Until* language over both frame classes.

1 Introduction

We examine the computational complexity of satisfiability for hybrid logics over transitive frames. This is an important frame class, because transitivity is the minimal requirement in many temporal applications, for example temporal verification. Modal, hybrid, and first-order logic over transitive models have been studied recently in [3, 14, 31, 19, 20, 18, 11]. Although the complexity of hybrid (tense) logics has been extensively examined [7, 15, 2, 3, 13], there are highly expressive hybrid languages for whose satisfiability problems only results over arbitrary, but not over restricted, temporally relevant, frame classes have been known.

Hybrid Languages are extensions of the language of modal logic that allow for naming and accessing states of a model explicitly. This renders hybrid logic an adequate representation formalism for many applications where the basic modal and/or temporal language does not suffice. Moreover, reasoning systems are easier to devise for hybrid than for modal logic.

Hybrid logic, as well as the foundations of temporal logic, goes back to Arthur Prior [26]. Since then, many — more or less powerful — languages have been studied. Here we briefly introduce the extensions that shall concern us in this paper.

Nominals are special atomic formulae that name states of models. They allow, for instance, for an axiom expressing irreflexivity, which cannot be captured by modal formulae: $i \rightarrow \neg \Diamond i$.

The *at operator* @ can be used to directly jump to states named by nominals, independently of the accessibility relation. Hence, the above formula could also be written as $@_i \neg \Diamond i$.

With the help of the *downarrow operator* \downarrow , it is possible to bind variables to states. Whenever $\downarrow x$ is encountered during the evaluation of a formula, the variable x is bound to the actual state s . All occurrences of x in the scope of this \downarrow are treated like nominals naming s . As an example, the formula $\downarrow x. \neg \Diamond \Diamond x$ reads as: Name the actual state x and make sure that it is not possible to reach x in two steps. This is an axiom for asymmetry, another property not expressible in modal logic.

Combined with the @ operator, \downarrow leads to a very powerful language that can formulate many desirable properties and goes far beyond the scope of the simple nominal language. To give a more impressive example, we consider the *Until operator*. The formula $U(\varphi, \psi)$ reads as “there is a point in the future at which φ holds, and at all points between now and this point, ψ holds”. What the basic modal language is not able to express, can be achieved by the hybrid \downarrow -@ language.

$$U(\varphi, \psi) \equiv \downarrow x. \Diamond \downarrow y. \varphi \wedge @_x \Box (\Diamond y \rightarrow \psi).$$

Besides more advanced temporal concepts such as “until” and “since”, hybrid temporal languages can express other desirable temporal notions such as “now”, “yesterday”, “today”, or “tomorrow”. Moreover, with hybrid logic one can capture many temporally relevant frame properties (besides the above mentioned, antisymmetry, trichotomy, directedness, ...). For this reason, hybrid temporal languages are of great interest where basic temporal logic reaches its limits [9, 5, 4, 13].

Why Transitivity? Hybrid logic is interpreted over Kripke frames and models, as is modal logic. A frame consists of a set of states (points in time) and an accessibility relation (where xRy says that y is in the future of x). In most temporal applications one requires this relation to be transitive. We concentrate on transitive frames because transitivity is a property that the future relations of many different temporal applications have in common, even if they differ in other properties such as tree-likeness, trichotomy, irreflexivity, or asymmetry.

But there are other reasons why this frame class is of interest, particularly in connection with computational complexity. In the special case of linear frames, nominals and @ can be simulated using the conventional modal operator and its converse. They do not add expressive power to the language in this case. The \downarrow operator is useless even on transitive trees, a representation of branching time. Though the class of transitive frames (and transitive trees for \downarrow -free hybrid languages, respectively) is a restricted frame class, it is general enough to separate hybrid from modal languages in terms of expressive power.

Yet another reason for considering precisely transitive frames will become clear in the next paragraph.

Complexity of Hybrid Logics. We use the complexity classes NP, PSPACE, EXP, NEXP, $n\text{EXP}$, $n \geq 2$, and coRE as known from [25]. A problem is nonelementarily decidable if it is decidable and not contained in any $n\text{EXP}$.

It goes without saying that reasoning tasks for richer logics require more resources than those for simpler languages such as the basic modal language. We focus on one reasoning task, namely satisfiability. The modal and temporal satisfiability problems over arbitrary as well as transitive frames are PSPACE-complete [23, 30]. If the “somewhere” modality \mathbf{E} is added, satisfiability becomes EXP-complete over arbitrary frames [29]. For many, more restricted, frame classes, modal and temporal satisfiability is NP-complete [23, 24, 28]. In contrast, the known part of the complexity spectrum of hybrid satisfiability reaches up to undecidability.

Many complexity results for hybrid languages have been established in [2, 3]. It was shown in [2] that the hybrid language with nominals and $\textcircled{\bullet}$ has a PSPACE-complete satisfiability problem and that satisfiability for the hybrid tense language is EXP-complete, even if $\textcircled{\bullet}$ or \mathbf{E} are added. It was proven in [3] that these problems have the same complexity (or drop to PSPACE-complete or NP-complete, respectively) if the class of frames is restricted to transitive frames (or transitive trees, or linear frames, respectively).

Moreover, the authors of [3] established EXP-completeness of satisfiability for the hybrid Until/Since language. The complexity of this language over transitive frames and transitive trees, respectively, has been open. PSPACE-completeness over linear frames is known from [13]. We want to find out at which exact requirements to the frame classes the decrease from EXP to PSPACE takes place.

Undecidability results for languages containing \downarrow originate from [7, 15]. The strongest such result, namely for a restricted fragment of the \downarrow language, is given in [2]. In recent work [33], it was shown that decidability of the \downarrow language can be regained under certain restrictions on the frame classes. It is possible that transitivity is another property under which the \downarrow language can be “tamed”.

Moreover, we have already observed that over transitive trees and linear orders, the \downarrow operator on its own is useless. Hence satisfiability over these frame classes is the same as for modal logic, namely complete for PSPACE and NP, respectively. This makes it more likely that we can “tame” \downarrow over transitive frames. But if so, to what extent? What happens if we allow for interactions of \downarrow with $\textcircled{\bullet}$ or the backward modality?

New Road-Map Pages. This paper establishes two groups of complexity results for hybrid languages over transitive frames and transitive trees.

First we examine satisfiability of the hybrid \downarrow language. Our most surprising result is the “taming” of this language over transitive frames: the satisfiability problem is NEXP-complete. This high level of complexity is retained even over complete frames (clusters). We also show that enriching the language by the backward-looking modality \mathbf{P} or the $\textcircled{\bullet}$ operator leads to undecidability in the case of transitive frames. Over transitive trees, the situation is different. Decidability for even the richest \downarrow language is easy to see, but it is nonelementary, as we will show.

As a second step, we consider satisfiability over transitive frames and transitive trees for the hybrid Until/Since-E language. We establish EXP-hardness for not more than the modal language with Until only. This is matched by an EXP upper bound for the full language in the case of transitive trees. As for transitive frames, we give a 2EXP upper bound.

Table 1 gives an overview of the satisfiability problems considered in this paper (marked bold) and visualizes how our complexity results arrange into a collection of previously known results. It makes use of the denotation of hybrid languages introduced in Section 2. Complexity classes without addition stand for completeness results; “nonel.” stands for “nonelementarily decidable”. The work from which the results originate, is cited. Conclusions from surrounding results are abbreviated by “c.”. The question mark stands for an open question, but decidability follows from the last result in that column. As for $\mathcal{HL}_{U,S}^E$ over transitive frames and transitive trees, EXP-hardness even holds for \mathcal{ML}_U .

| hybrid lang. | complexity over arbitrary frames | complexity over transitive frames | complexity over transitive trees | complexity over linear orders |
|---|----------------------------------|--|----------------------------------|-------------------------------|
| \mathcal{HL}° | PSPACE [2] | PSPACE [3] | PSPACE [3, c.] | NP [3, c.] |
| $\mathcal{HL}_{F,P}$ | EXP [2] | EXP [3] | PSPACE [3] | NP [3] |
| $\mathcal{HL}_{F,P}^E$ | EXP [3] | EXP [3] | PSPACE [3] | NP [3] |
| $\mathcal{HL}_{U,S}^E$ | EXP [3] | in 2EXP (Th. 14), EXP-hard (Th. 12) | EXP (Th. 12,15) | PSPACE-hard [27] |
| \mathcal{HL}^\downarrow | coRE [2] | NEXP (Th. 1) | PSPACE [3] | NP [13] |
| $\mathcal{HL}^{\downarrow,\circ}$ | coRE [2] | coRE (Th. 10) | nonel. (Th. 11) | ? |
| $\mathcal{HL}_{F,P}^\downarrow$ | coRE [2, c.] | coRE (Th. 10) | nonel. (Th. 11) | nonel. [13, c.] |
| $\mathcal{HL}_{F,P}^{\downarrow,\circ}$ | coRE [2, c.] | coRE (c.) | nonel. (Th. 11) | nonel. [13] |

Table 1. An overview of complexity results for hybrid logics.

Legend. This paper is organized as follows. In Section 2, we give all necessary definitions and notations of modal and hybrid logic as well as an overview of our results in the context of previous work. We present the decidability and undecidability results for the hybrid \downarrow languages in Sections 3 and 4. The hybrid Until/Since language is examined in section 5, and Section 6 ends the paper with some concluding remarks.

2 Modal and Hybrid Logic

We define the basic concepts and notations of modal and hybrid logic that are relevant for this paper. The fundamentals of modal logic can be found in [6]; those of hybrid logic in [2, 5].

Modal Logic. Let PROP be a countable set of *propositional atoms*. The language \mathcal{ML} of modal logic is the set of all formulae of the form $\varphi ::= p \mid \neg\varphi \mid$

$\varphi \wedge \varphi' \mid \Diamond \varphi$, where $p \in \text{PROP}$. We use the well-known abbreviations \vee , \rightarrow , \leftrightarrow , \top (“true”), and \perp (“false”), as well as $\Box \varphi := \neg \Diamond \neg \varphi$.

A (*Kripke*) *model* is a triple $\mathcal{M} = (M, R, V)$, where M is a nonempty set of *states*, $R \subseteq M \times M$ is a binary relation — the *accessibility relation* —, and $V : \text{PROP} \rightarrow \mathfrak{P}(M)$ is a function — the *valuation function*. The structure $\mathcal{F} = (M, R)$ is called a *frame*. Given a model $\mathcal{M} = (M, R, V)$ and a state $m \in M$, the *satisfaction relation* is defined by

$$\begin{aligned} \mathcal{M}, m \models p & \quad \text{iff } m \in V(p), \ p \in \text{PROP}, \\ \mathcal{M}, m \models \neg \varphi & \quad \text{iff } \mathcal{M}, m \not\models \varphi, \\ \mathcal{M}, m \models \varphi \wedge \psi & \quad \text{iff } \mathcal{M}, m \models \varphi \ \& \ \mathcal{M}, m \models \psi, \\ \mathcal{M}, m \models \Diamond \psi & \quad \text{iff } \exists n \in M (mRn \ \& \ \mathcal{M}, n \models \psi). \end{aligned} \quad (2.1)$$

A formula φ is *satisfiable* if there exist a model $\mathcal{M} = (M, R, V)$ and a state $m \in M$, such that $\mathcal{M}, m \models \varphi$. If all states from \mathcal{M} satisfy φ , we write $\mathcal{M} \models \varphi$ and say that φ is *globally satisfied* by \mathcal{M} .

Temporal Logic. The language of temporal logic (tense logic) is the set of all formulae of the form $\varphi ::= p \mid \neg \varphi \mid \varphi \wedge \varphi' \mid F\varphi \mid P\varphi$, where $p \in \text{PROP}$. It is common practice to use the abbreviations $G\varphi := \neg F\neg \varphi$ and $H\varphi := \neg P\neg \varphi$. The operator F replaces \Diamond , hence satisfaction for F -formulae is defined as in (2.1). In the case of P -formulae, the term mRn must be replaced by nRm .

Whenever one wants to speak not only of states accessible from the actual state, but also of states “between” the actual and some accessible state, one can make use of the binary operators U (“until”) and S (“since”), for which satisfaction is defined by

$$\begin{aligned} \mathcal{M}, m \models U(\varphi, \psi) & \quad \text{iff } \exists n \in M (mRn \ \& \ \mathcal{M}, n \models \varphi \ \& \ \forall s \in M (mRsRn \Rightarrow \mathcal{M}, s \models \psi)), \\ \mathcal{M}, m \models S(\varphi, \psi) & \quad \text{iff } \exists n \in M (nRm \ \& \ \mathcal{M}, n \models \varphi \ \& \ \forall s \in M (nRsRm \Rightarrow \mathcal{M}, s \models \psi)). \end{aligned}$$

The U/S language is strictly stronger than the basic temporal language in the sense that F and P can be expressed by U and S (e.g. $F\varphi = U(\varphi, \top)$), but not vice versa.

In [3], a variant of the U/S operators, U^+ and S^+ , is introduced. Satisfaction for U^+ (analogously for S^+) is defined by

$$\begin{aligned} \mathcal{M}, m \models U^+(\varphi, \psi) & \quad (2.2) \\ \text{iff } \exists n \in M (mRn \ \& \ \mathcal{M}, n \models \varphi \ \& \ \forall s \in M (mR^+sR^+n \Rightarrow \mathcal{M}, s \models \psi)), \end{aligned}$$

where R^+ is the transitive closure of R . With the help of U^+ and S^+ , the authors of [3] “simulated” transitive frames syntactically. We go a step further and define another modification, U^{++} and S^{++} , with the satisfaction relation from (2.2), where the last remaining term mRn is replaced by mR^+n . This temporal language is an even closer simulation of transitivity, as we will see in Section 5.

Hybrid Logic. As indicated in the previous section, *the* hybrid language does not exist. Rather there are several extensions of the modal language allowing

for explicit references to states and therefore being called hybrid. We introduce the richest of those hybrid languages that will interest us in this paper. The definitions and notations are taken from [2, 3].

Let NOM be a countable set of *nominals*, SVAR be a countable set of *state variables*, and $\text{ATOM} = \text{PROP} \cup \text{NOM} \cup \text{SVAR}$. It is common practice to write propositional atoms as p, q, \dots , nominals as i, j, \dots , and state variables as x, y, \dots . The *full hybrid language* $\mathcal{HL}^{\downarrow, @}$ is the set of all formulae of the form $\varphi ::= a \mid \neg\varphi \mid \varphi \wedge \varphi' \mid \Diamond\varphi \mid @_t\varphi \mid \downarrow x.\varphi$, where $a \in \text{ATOM}$, $t \in \text{NOM} \cup \text{SVAR}$, and $x \in \text{SVAR}$.

A formula is called *pure* iff it contains no propositional atoms; *nominal-free* iff it contains no nominals; and a *sentence* iff it contains no free state variables. (*Free* and *bound* are defined as usual; the only binding operator here is \downarrow .)

A *hybrid model* is a Kripke model with the valuation function V extended to $\text{PROP} \cup \text{NOM}$, where for all $i \in \text{NOM}$, $|V(i)| = 1$. Whenever it is clear from the context, we will omit the word “hybrid” when referring to models. In order to evaluate \downarrow -formulae, an *assignment* $g : \text{SVAR} \rightarrow M$ for \mathcal{M} is necessary. Given an assignment g , a state variable x and a state m , an *x -variant* g_m^x of g is defined by $g_m^x(x) = m$ and $g_m^x(x') = g(x')$ for all $x' \neq x$. For any atom a , let $[V, g](a) = \{g(a)\}$ if $a \in \text{SVAR}$, and $V(a)$, otherwise. The satisfaction relation for hybrid formulae is defined by

$$\begin{aligned} \mathcal{M}, g, m \models a & \quad \text{iff } m \in [V, g](a), \ a \in \text{ATOM}, \\ \mathcal{M}, g, m \models \neg\varphi & \quad \text{iff } \mathcal{M}, g, m \not\models \varphi, \\ \mathcal{M}, g, m \models \varphi \wedge \psi & \quad \text{iff } \mathcal{M}, g, m \models \varphi \ \& \ \mathcal{M}, g, m \models \psi, \\ \mathcal{M}, g, m \models \Diamond\varphi & \quad \text{iff } \exists n \in M(mRn \ \& \ \mathcal{M}, g, n \models \varphi), \\ \mathcal{M}, g, m \models @_t\varphi & \quad \text{iff } \mathcal{M}, g, n \models \varphi \ \& \ [V, g](t) = \{n\}, \\ \mathcal{M}, g, m \models \downarrow x.\varphi & \quad \text{iff } \mathcal{M}, g_m^x, m \models \varphi. \end{aligned}$$

A formula is *satisfiable* if there exist a model $\mathcal{M} = (M, R, V)$, an assignment g for \mathcal{M} , and a state $m \in M$, such that $\mathcal{M}, g, m \models \varphi$.

We sometimes use the “somewhere” modality \mathbf{E} having the interpretation $\mathcal{M}, g, m \models \mathbf{E}\varphi$ iff $\exists n \in M(\mathcal{M}, g, n \models \varphi)$. In this case, $@$ is needless, because $@_t\varphi$ can be expressed by $\mathbf{E}(t \wedge \varphi)$.

First-order Logic. Modal and hybrid logic can be embedded into fragments of first-order logic (FOL). We will always use the standard notation of FOL.

We will make use of certain fragments of FOL and denote them in the style of [10]: $[\text{all}, (u, 1)]$, where $u \in \omega$. This notation stands for the fragment without equality, without function symbols, and with no other relation symbols than one binary and u unary ones. We denote the satisfiability problem for such a fragment by $[\text{all}, (u, 1)]\text{-SAT}$ and $[\text{all}, (u, 1)]\text{-trans-SAT}$, where the latter requires that the binary relation symbol is interpreted by a transitive relation.

The *Standard Translation* ST [33] embeds hybrid logic into FOL and consists of two functions ST_x and ST_y defined recursively. Since ST_y is obtained from ST_x by exchanging x and y , we only give ST_x here.

$$\begin{aligned} ST_x(p) &= P(x), & ST_x(\Diamond\varphi) &= \exists y(xRy \wedge ST_y(\varphi)), \\ ST_x(t) &= t=x, & ST_x(@_t\varphi) &= \exists y(y=t \wedge ST_y(\varphi)), \\ ST_x(\neg\varphi) &= \neg ST_x(\varphi), & ST_x(\downarrow v.\varphi) &= \exists v(x=v \wedge ST_x(\varphi)), \\ ST_x(\varphi \wedge \psi) &= ST_x(\varphi) \wedge ST_x(\psi), & ST_x(E\varphi) &= \exists y(ST_y(\varphi)), \end{aligned}$$

where $p \in \text{PROP}$, $t \in \text{NOM} \cup \text{SVAR}$, and $v \in \text{SVAR}$.

Properties of Models and Frames. Let $\mathcal{M} = (M, R, V)$ be a (Kripke or hybrid) model with the underlying frame $\mathcal{F} = (M, R)$. By R^+ we denote the transitive closure of R . For any subset $M' \subseteq M$, we write $R|_{M'}$ and $V|_{M'}$ for the restrictions of R and V to M' . We will refer to *transitive frames* or *linear frames* whenever we mean frames whose accessibility relation is transitive or a linear order, respectively. A *linear order* is an irreflexive, transitive, and trichotomous ($\forall xy(xRy \text{ or } x=y \text{ or } yRx)$) relation. The frame \mathcal{F} is a *tree* iff it is acyclic and connected, and every point has at most one R -predecessor. A *transitive tree* is any (M, R^+) , where (M, R) is a tree.

Satisfiability Problems. Whenever we leave one or more operators out of the hybrid language, we omit the according superscript of \mathcal{HL} . If we proceed to a hybrid tense language, we add the suitable temporal operator(s) as subscript(s) to \mathcal{HL} . Analogously, when equipping the modal language with additional operators, we add them as sub- or superscripts to \mathcal{ML} .

For any hybrid language \mathcal{HL}_y^x , the *satisfiability problem* $\mathcal{HL}_y^x\text{-SAT}$ is defined as follows: Given a formula $\varphi \in \mathcal{HL}_y^x$, do there exist a hybrid model \mathcal{M} , an assignment g for \mathcal{M} , and a state $m \in M$ such that $\mathcal{M}, g, m \models \varphi$? If \downarrow is not in the considered language, then the assignment g may be left out of this formulation. If we only ask for *transitive models* (or *transitive trees* or *linear models*, respectively) satisfying φ , then we speak of $\mathcal{HL}_y^x\text{-trans-SAT}$ (or $\mathcal{HL}_y^x\text{-tt-SAT}$, or $\mathcal{HL}_y^x\text{-lin-SAT}$, respectively). As an example for our notation, the satisfiability problem over transitive frames for the hybrid temporal \downarrow language is denoted by $\mathcal{HL}_{F,P}^\downarrow\text{-trans-SAT}$.

3 Deciding \mathcal{HL}^\downarrow over Transitive Frames

In [2] Areces, Blackburn, and Marx proved that the downarrow operator \downarrow turns the satisfiability problem for hybrid logics undecidable in general, even if no interaction with $@$ or P is allowed. We prove that undecidability vanishes if frames are required to be transitive.

Theorem 1 *The satisfiability problem for \mathcal{HL}^\downarrow over transitive frames is complete for NEXP.*

Before we start with the proof we have a first look at \mathcal{HL}^\downarrow over transitive frames. Obviously, it has no finite model property. E.g., the following sentence requires a model containing an infinite chain of states labeled p .

$$p \wedge \Diamond p \wedge \Box \Diamond p \wedge \Box \downarrow x. \neg \Diamond x$$

Neither is it always possible to find a model that is a transitive tree. But, in some way, we can get close to this. Although our models may contain cycles, transitivity ensures that all states in a cycle are pairwise connected. I.e., the subframe consisting of these states is complete. Therefore, we can view a model as consisting of maximal complete subframes and single states, that are connected in a transitive but acyclic fashion.

For every transitive model $\mathcal{M} = (M, R, V)$ we define $B(\mathcal{M}) = (M', R', V')$ to be the model defined as follows. First, we replace each maximal complete subframe with a single vertex. Second, we unravel the resulting structure into a (potentially infinite) transitive tree T . Then we replace each vertex of T by (a copy of) the clique of \mathcal{M} from which it is derived. We will refer to the resulting model as the *block tree* of \mathcal{M} . Note that a block tree is not a tree but its clique frame is a transitive tree.

In the following we are often interested in the underlying tree structure of a block tree and refer to the cliques of a block tree as *nodes*. For a state s , we denote its node by u_s . We say that a node v is *below* a node u if the states of v are reachable from the states in u (but not vice versa) and a node v is a *child* of a node u , if v is below u but there is no node w below u and above v .

Likewise we use the terms tree, subtree, and leaf for block tree, sub block tree and leaf clique, respectively.

We have to be careful about how to treat nominals when unraveling a model. If $V(i) = \{s\}$ for a nominal i and a state $s \in M$, we define $V'(i)$ to be the set of states from M' , that are copies via the unraveling of s . Therefore, $B(\mathcal{M})$ is not a model as defined in Section 2 because nominals may hold at more than one state, but by viewing i as a propositional atom it can be treated as a model. The satisfaction relation is not affected and it is therefore easy to see that this transformation preserves satisfaction of \mathcal{HL}^\downarrow -sentences: The relation associating each state of \mathcal{M} with every copy in $B(\mathcal{M})$ is a quasi-injective bisimulation³ [8].

Lemma 2 *For every transitive model \mathcal{M} and every \mathcal{HL}^\downarrow -sentence φ :*

$$\mathcal{M} \models \varphi \iff B(\mathcal{M}) \models \varphi.$$

Note that we can always get a model for φ from $B(\mathcal{M})$ by joining the states labeled with the same nominals, but this model might be different from \mathcal{M} .

Before we show how to use this tree-like structure to decide \mathcal{HL}^\downarrow over transitive frames, we focus on complete subframes and show that their size can be bounded.

³ The notion of bisimulation has to be extended by requiring states carrying the same nominal to be related.

3.1 $\mathcal{H}\mathcal{L}^\downarrow$ over Complete Frames

As complete subframes are a significant part of transitive models for $\mathcal{H}\mathcal{L}^\downarrow$ -sentences, we are now going to study the satisfiability problem of $\mathcal{H}\mathcal{L}^\downarrow$ over complete frames. The most important result for our purpose is an exponential-size model property of $\mathcal{H}\mathcal{L}^\downarrow$ over complete frames.

We want to start by giving some insight why this property holds. In complete frames, the accessibility relation does not distinguish different states. Of course, states can be told apart if they are labeled differently by propositions. But the number of different labelings is exponentially bounded in the size of the formula. To use more states, we have to distinguish states labeled equally. This can only be done by assigning names to these states. But the number of states we can distinguish in this way is bounded by the number of different state variables.

While intuition is clear, we can prove this bound by observing that $\mathcal{H}\mathcal{L}^\downarrow$ over complete frames is equivalent to the *Monadic Class with equality* ($\text{MC}_=$), the fragment of first-order logic with only unary predicates, equality, and no function symbols [10].

This result allows us to transfer complexity results and model properties for $\text{MC}_=$ [10] to $\mathcal{H}\mathcal{L}^\downarrow$ over complete frames.

Theorem 3 *$\mathcal{H}\mathcal{L}^\downarrow$ over complete frames has the exponential-size model property and its satisfiability problem is complete for NEXP.*

The lower bound can be transferred directly to the case of transitive frames.

Corollary 4 *The satisfiability problem for $\mathcal{H}\mathcal{L}^\downarrow$ over transitive frames is hard for NEXP.*

3.2 On Transitive Frames for $\mathcal{H}\mathcal{L}^\downarrow$

Let us summarize what we have seen so far. For every $\mathcal{H}\mathcal{L}^\downarrow$ -formula satisfiable over transitive frames, instead of a transitive model we can consider its block tree. The size of the cliques in the block tree can be exponentially bounded in the size of the formula by Corollary 3.

The algorithm for testing $\mathcal{H}\mathcal{L}^\downarrow$ -satisfiability will essentially guess a model and verify that it is correct. As there is no finite model property, all models might be infinite. Nevertheless, we will show that, if the formula is satisfiable, there is always a model with a regular structure in which certain finite patterns are repeated infinitely often. This will allow us to find a finite representation of such a model.

To this end, Definition 5 captures the information about a state of a block tree that will be needed for the following. Intuitively, the φ -type of a state captures the information needed about its subtree in order to evaluate a subformula of a given formula φ at this state. Here, $\psi[free/\perp]$ is the sentence obtained from ψ by replacing every free variable by \perp and $sub(\varphi)$ is the set of all subformulae of φ .

Definition 5 Let φ be a \mathcal{HL}^\downarrow -sentence and $\mathcal{B} = (M, R, V)$ a block tree model which is a model of φ . The φ -type of a state $s \in M$ is the set of all sentences from $\{\psi[\text{free}/\perp] \mid \Diamond\psi \in \text{sub}(\varphi)\}$ that hold at some state in the subtree rooted at s .

Note that states in the same clique have the same φ -type. Therefore, we can speak of the φ -type of a node. The type of a node is always a superset of the types of its children. More precisely, it is always the union of the types of the children together with the set of relevant formulae which hold in the node itself.

When evaluating a subformula of a \mathcal{HL}^\downarrow -sentence φ at some state s of a block tree, all we need to know about states strictly below u_s are the φ -types of the children of u_s . I.e., we can replace subtrees below u_s by subtrees of the same φ -type. In the following lemma, for a block tree \mathcal{B} and two states s_1, s_2 , $\mathcal{B}[s_1/s_2]$ denotes the block tree resulting from \mathcal{B} by replacing the subtree rooted at s_1 by the subtree rooted at s_2 . The result of this substitution is again a block tree.

Lemma 6 Let φ be a \mathcal{HL}^\downarrow -sentence, $\mathcal{B} = (M, R, V)$ a block tree model of φ and s_1 and s_2 states of \mathcal{M} such that there is a path from s_1 to s_2 but not vice versa. For every formula $\psi \in \text{sub}(\varphi)$, every state s_3 of \mathcal{M} of the same φ -type as s_2 , and every assignment g that maps all free variables in ψ to states in M preceding s_1 :

$$\mathcal{B}, g, s_1 \models \psi \iff \mathcal{B}[u_{s_2}/u_{s_3}], g, s_1 \models \psi.$$

Note that we restricted the choice of g only to those assignments that are really relevant when evaluating the sentence φ .

We can use the previous lemma to get some nice restrictions on the block trees under consideration. E.g., we can assume that for every sentence in the type of a node, there is a witness in the node itself or in one of its children.

Lemma 7 Let φ be a \mathcal{HL}^\downarrow -sentence satisfiable over transitive frames. Then there is a block tree model \mathcal{B} for φ , in which

- every node has at most $|\varphi|$ children,
- for every node u with φ -type t and every \mathcal{HL}^\downarrow -sentence $\psi \in t$, ψ holds at a state in u or at a state in a child of u , and
- on every path from the root, infinite or ending at a leaf, every φ -type occurs only once or infinitely often.

3.3 Deciding \mathcal{HL}^\downarrow -SAT over Transitive Frames

We will now finish the proof of Theorem 1 by presenting a nondeterministic algorithm that decides \mathcal{HL}^\downarrow -SAT over transitive frames in exponential time, basically by guessing and verifying the finite representation of a block tree model for a given \mathcal{HL}^\downarrow -sentence φ .

Given a block tree \mathcal{B} with the properties of Lemma 7, we get a finite representation as follows. For each path of \mathcal{B} we consider the first node v that has the same type as its parent node u . We replace the subtree below v by a single state labeled with a reference to u . We need to keep v because it might be the only

witness for a formula in the φ -type of u (cf. Lemma 7). Clearly, the resulting structure is finite.

By Lemma 6 and Lemma 7, we can get a block tree model from this representation by replacing each reference with the subtree rooted at the referenced node, i.e., essentially by an unraveling.

Due to Lemma 6, the size of the representation can be reduced even further. If there are two nodes u and v of the same φ -type which are not on the same path and both are the first node of their type on their path from the root, we can replace the subtree rooted at v with the subtree of u . I.e., whenever two nodes have the same φ -type, we can assume that their generated subtrees are equal. We have to check them only once.

Such a representation can be described by a structure $(M \dot{\cup} C, R, V)$ such that the states in C have no outgoing edges, and a function f from C to M . Note that a state in C is a node of its own, in fact a leaf, and cannot be in the same complete subframe as a state in M . A state $s \in C$ stands for a repetition respectively duplication of the subtree rooted at $f(s)$, including states from C . This causes infinite repetition if s is below $f(s)$.

This observation can be reflected in our representation by replacing every duplicate with a reference. This causes every type to appear at most twice in the representation, thus the number of nodes is at most exponential.

Summing up, if φ is satisfiable, there is a representation for of a block tree model for φ of size at most exponential in the length of φ . The first step of the algorithm is to guess such a representation (step 1).

In order to obtain an algorithm which tests whether the representations indeed represents a model of φ , we describe how to modify the model checking algorithm MCFULL by Franceschet and de Rijke presented in [12] to do so. First, we deal with the states in C (step 2). Next, the model checking algorithm MCFULL is used on the states in M (step 3). We have to modify this algorithm in two respects. First, it has to use the information guessed for the states in C . Second, it should compute the φ -types of the states in M . To this end, it first evaluates the sentences resulting from subformulae of φ by replacing free variables with \perp . We call this modified algorithm MCFULL'. The changes are straightforward.

After running MCFULL' the algorithm has computed for each state the set of formulae that hold at this state. These sets depend on the guesses in Step 2. Therefore, the algorithm has to verify the consistency of these guesses. This can be done by comparing the φ -types of the states in C with the types of the referenced states (step 4).

The last two steps are easy. The algorithm checks that the φ -types are consistent and reject if this is not the case (step 5). Finally, it checks if φ holds at some state in M (step 6).

Our algorithm for $\mathcal{H}\mathcal{L}^\downarrow$ -satisfiability over transitive frames

- 1: Guess the representation of a block tree.
- 2: Guess a φ -type for every state in C .
- 3: Run MCFULL' for the states in M .
- 4: Compute the φ -type for every state in M referenced by a state in C .
- 5: Check for every state $s \in C$: $f(s)$ has the same φ -type as s . If not, reject.
- 6: Accept iff φ holds at some state in M .

Theorem 8 *The algorithm presented above decides $\mathcal{H}\mathcal{L}^\downarrow$ -satisfiability over transitive frames nondeterministically in exponential time.*

Proof. In Section 3.2, we have seen that for every satisfiable sentence φ there is a block tree model as described in Lemma 7. We have also seen how to represent this block tree in a finite manner. The algorithm can guess this representation and the φ -types of the states in C . The computation of the φ -types of the states in M works correctly, because we have a witness for every sentence in the type in our representation. This is by Lemma 7, which ensures that witnesses are in the node of the state or in one of its children. Therefore, we cut below these witnesses when building the finite representation. Consequently, the algorithm will accept.

On the other hand, if the algorithm accepts, it is straightforward to construct a block tree model from the guessed representation. The only critical point for soundness is the verification of the φ -types guessed in Step 2, more precisely, the computation of the φ -types of the states in M . First, the φ -type of some state $s \in M$ contains only sentences that hold at some state of the represented model below s . This can be assured by looking only at states in M and not at states in C . That the φ -type of s contains all sentences that hold below s can be assured by following the links represented by states in C .

The first two steps of the algorithm can be performed in exponential time because the representation is of at most exponential size. That Step 3 runs in exponential time follows from Theorem 4.5 of [12], the truth of which is not affected by our modifications. The time bounds for the other steps follow again from the exponential size bound of the representation. \square

From Theorem 8 and Corollary 4 we can directly conclude Theorem 1.

4 Richer hybrid \downarrow Logics over Transitive Frames and Transitive Trees

This section is concerned with satisfiability over transitive frames and transitive trees for extensions of $\mathcal{H}\mathcal{L}^\downarrow$.

Our first point is that we cannot sustain decidability over transitive frames if we enrich $\mathcal{H}\mathcal{L}^\downarrow$ with @ or the backward looking modality P. We prove undecidability in both cases, proceeding in two steps. First, we show that [all, (4, 1)]-trans-SAT is undecidable. This is done by a reduction from [all, (0, 1)]-SAT. The undecidability of the latter is a consequence of the undecidability of contained

traditional standard classes [10]. The second step consists of reduction from $[\text{all}, (4, 1)]$ -trans-SAT to $\mathcal{HL}^{\downarrow, \odot}$ -trans-SAT and $\mathcal{HL}_{F, P}^{\downarrow}$ -trans-SAT, respectively. To be more precise, the ranges of these reductions will be the fragments of the respective hybrid languages consisting of all nominal-free sentences.

Lemma 9 $[\text{all}, (4, 1)]$ -trans-SAT is undecidable.

Proof. In order to obtain the required reduction from $[\text{all}, (0, 1)]$ -SAT, we will transform a (not necessarily transitive) model satisfying α into a transitive one. Simply taking the transitive closure in most cases adds new pairs to the interpretation of the relation and is not sufficient for keeping the information which pairs were in the “old” relation and which pairs were not. This problem does not arise if we instead use a variation of the *zig-zag technique* successfully applied in [3] for a reduction between a modal and a hybrid language. The core idea of this technique is to simulate an R -step $t_1 R t_2$ in the original model $\mathcal{M} = (D, I)$ by a zig-zag transition in a model $\mathcal{M}' = (D', I')$, where $I'(R)$ is transitive, as shown in Figure 1.

We define a translation function $(\cdot)^t$ using four extra predicate symbols $0, 1, 2, 3$ as follows.




Fig. 1. A zig-zag transition.

The translation of the xRy -atoms exactly reflects the shown zig-zag transition. It is straightforward to prove the following claim: *For each formula α , α is satisfiable iff $f(\alpha)$ is satisfiable in some model that interprets R by a transitive relation.* Since $(\cdot)^t$ is an appropriate (even polynomial-time) reduction function, we have established undecidability for $[\text{all}, (4, 1)]$ -trans-SAT. \square

Theorem 10 $\mathcal{HL}^{\downarrow, \odot}$ -trans-SAT and $\mathcal{HL}_{F, P}^{\downarrow}$ -trans-SAT are undecidable.

Proof. We reduce $[\text{all}, (4, 1)]$ -trans-SAT to $\mathcal{HL}^{\downarrow, \odot}$ -trans-SAT and $\mathcal{HL}_{F, P}^{\downarrow}$ -trans-SAT, respectively, invoking a spypoint argument (see [7, 2]). A spypoint is a state s of a hybrid model that sees all other states and is named by a fresh nominal i . Since our reduction will not make use of any nominals, we can establish this undecidability result for the nominal-free fragments of the hybrid languages in question. We simply treat i as a state variable and bind it to s .

We first treat the case of $\mathcal{HL}^{\downarrow, \odot}$ and define a translation function $(\cdot)^t$ from the first-order fragment to $\mathcal{HL}^{\downarrow, \odot}$ by $(xRy)^t = @_x \odot y$, $(P(x))^t = @_x p$, $(\neg\alpha)^t = \neg(\alpha^t)$, $(\alpha \wedge \beta)^t = \alpha^t \wedge \beta^t$, and $(\exists x \alpha)^t = @_i \odot \downarrow x. \alpha^t$. The (polynomial) reduction function f is defined by $f(\alpha) = \downarrow i. (\neg \odot i \wedge \odot \alpha^t)$.

It is straightforward to show that each formula α is satisfiable iff $f(\alpha)$ is satisfiable, and to adopt the translation function $(\cdot)^t$ to the case of $\mathcal{H}\mathcal{L}_{F,P}^\downarrow$. \square

The second point of this section is that over transitive trees, where decidability of $\mathcal{H}\mathcal{L}^\downarrow$ is trivial, even the extension $\mathcal{H}\mathcal{L}_{F,P}^{\downarrow,\oplus}$ is decidable. This is an immediate consequence of the decidability of the monadic second-order theory of the countably branching tree, $\text{S}\omega\text{S}$, [10]. However, we have to face a nonelementary lower bound in both cases $\mathcal{H}\mathcal{L}^{\downarrow,\oplus}$ and $\mathcal{H}\mathcal{L}_{F,P}^\downarrow$, which is obtained by a reduction from the nonelementarily decidable $\mathcal{H}\mathcal{L}_{F,P}^\downarrow\text{-lin-SAT}$ [13].

Theorem 11 $\mathcal{H}\mathcal{L}_{F,P}^\downarrow\text{-tt-SAT}$, $\mathcal{H}\mathcal{L}^{\downarrow,\oplus}\text{-tt-SAT}$, and $\mathcal{H}\mathcal{L}_{F,P}^{\downarrow,\oplus}\text{-tt-SAT}$ are nonelementary decidable.

5 Hybrid Until/Since Logic over Transitive Frames and Transitive Trees

In this section, we will consider $\mathcal{H}\mathcal{L}_{U,S}^E\text{-trans-SAT}$ and $\mathcal{H}\mathcal{L}_{U,S}^E\text{-tt-SAT}$. In [3] it was shown that $\mathcal{H}\mathcal{L}_{U,S}^{\oplus}\text{-SAT}$ is EXP-complete.

As for the lower bound, we establish a result as general as possible, namely EXP-hardness of $\mathcal{M}\mathcal{L}_U\text{-trans-SAT}$ and $\mathcal{M}\mathcal{L}_U\text{-tt-SAT}$.

Theorem 12 $\mathcal{M}\mathcal{L}_U\text{-trans-SAT}$ and $\mathcal{M}\mathcal{L}_U\text{-tt-SAT}$ are EXP-hard.

Proof. We will reduce the *global* satisfiability problem for $\mathcal{M}\mathcal{L}$ to both our (local) problems $\mathcal{M}\mathcal{L}_U\text{-trans-SAT}$ and $\mathcal{M}\mathcal{L}_U\text{-tt-SAT}$ using the same reduction function. The global satisfiability problem is defined by $\mathcal{M}\mathcal{L}\text{-GLOBSAT} = \{\varphi \in \mathcal{M}\mathcal{L} \mid \varphi \text{ is true in all states of some Kripke model } \mathcal{M}\}$. Its EXP-completeness is a direct consequence of the EXP-completeness of $\mathcal{M}\mathcal{L}^E\text{-SAT}$ [29].

It may seem difficult to try reducing this problem over *arbitrary* frames to our satisfiability problem over *transitive* frames. The critical point lies in making a non-transitive model transitive: taking the transitive closure of its relation forces us to add new accessibilities that would disturb satisfaction of $\neg\Diamond$ -formulae. Fortunately though, the U operator can make us distinguish the accessibilities from the original model from those that have been added to make the relation transitive. Hence, a translation of $\Diamond\varphi$ should demand: “Make sure that the actual state sees a state in which the translation of φ holds, and that there is no state in between.” This translates as $U(\varphi^t, \perp)$ into the modal language.

To construct the required reduction, we define a translation function $(\cdot)^t : \mathcal{M}\mathcal{L} \rightarrow \mathcal{M}\mathcal{L}_U$ by $p^t = p$, $p \in \text{PROP}$; $(\varphi \wedge \psi)^t = \varphi^t \wedge \psi^t$; $(\neg\varphi)^t = \neg(\varphi^t)$; and $(\Diamond\varphi)^t = U(\varphi^t, \perp)$. Using $(\cdot)^t$, we construct a reduction function $f : \mathcal{M}\mathcal{L} \rightarrow \mathcal{M}\mathcal{L}_U$ — that is clearly computable in polynomial time — via $f(\varphi) = \varphi^t \wedge \Box\varphi^t$. It is straightforward to prove the following two claims for each $\varphi \in \mathcal{M}\mathcal{L}$.

- (1) If $\varphi \in \mathcal{M}\mathcal{L}\text{-GLOBSAT}$, then $f(\varphi) \in \mathcal{M}\mathcal{L}_U\text{-tt-SAT}$.
- (2) If $f(\varphi) \in \mathcal{M}\mathcal{L}_U\text{-trans-SAT}$, then $\varphi \in \mathcal{M}\mathcal{L}\text{-GLOBSAT}$. \square

The upper bounds for $\mathcal{HL}_{U,S}^E$ -trans-SAT and $\mathcal{HL}_{U,S}^E$ -tt-SAT require separate treatment. As for $\mathcal{HL}_{U,S}^E$ -trans-SAT, we use an embedding into an appropriate fragment of first-order logic. In order to eliminate transitivity, we “simulate” this semantic property by syntactic means, namely using the operators U^{++} and S^{++} defined in Section 2.

Lemma 13 *For any $X \subseteq \{ @, E \}$, $\mathcal{HL}_{U,S}^X$ -trans-SAT and $\mathcal{HL}_{U^{++},S^{++}}^X$ -SAT are polynomially reducible to each other.*

Now it is not difficult anymore to obtain a 2EXP upper bound for $\mathcal{HL}_{U,S}^@$ -trans-SAT by an embedding into the loosely μ -guarded fragment μLGF of first-order logic whose satisfiability problem is 2EXP-complete [17]. Only the E operator requires a more careful analysis.

Theorem 14 *$\mathcal{HL}_{U,S}^E$ -trans-SAT is in 2EXP.*

To show that $\mathcal{HL}_{U,S}^E$ -tt-SAT is in EXP, too, we use an embedding into $\mathcal{PDL}_{\text{tree}}$, the propositional dynamic logic for sibling-ordered trees introduced in [21, 22]. Finite, node-labelled, sibling-ordered trees are the logical abstraction of XML (eXtensible Markup Language) documents. In [1], it was shown that satisfiability of $\mathcal{PDL}_{\text{tree}}$ formulae at the root of finite trees ($\mathcal{PDL}_{\text{tree}}$ -SAT) is decidable in EXP.

Theorem 15 *$\mathcal{HL}_{U,S}^E$ -tt-SAT is in EXP.*

6 Conclusion

We have established two groups of complexity results for hybrid logics over two temporally relevant frame classes, i.e. that of transitive frames and that of transitive trees.

First we have “tamed” \mathcal{HL}^\downarrow over transitive frames showing that \mathcal{HL}^\downarrow -trans-SAT is NEXP-complete. We have shown that, in contrast, $\mathcal{HL}^{\downarrow,@}$ -trans-SAT and $\mathcal{HL}_{F,P}^\downarrow$ -trans-SAT are undecidable. Over transitive trees, three enrichments of \mathcal{HL}^\downarrow are decidable, albeit nonelementarily, namely $\mathcal{HL}^{\downarrow,@}$ -tt-SAT, $\mathcal{HL}_{F,P}^\downarrow$ -tt-SAT, and $\mathcal{HL}_{F,P}^{\downarrow,@}$ -tt-SAT.

In the second part of our work, we established an EXP lower bound for \mathcal{ML}_U -trans-SAT and \mathcal{ML}_U -tt-SAT and matched the latter with an EXP upper bound for $\mathcal{HL}_{U,S}^E$ -tt-SAT. This is the same complexity as for satisfiability over *arbitrary* frames for the same language. As for $\mathcal{HL}_{U,S}^E$ -trans-SAT, we have given a 2EXP upper bound. We conjecture EXP-completeness.

Over linear frames, the complexity of hybrid U/S logic is still open. As a special case, satisfiability of $\mathcal{HL}_{U,S}^@$ over $(\mathbb{N}, >)$ is known to be PSPACE-complete [13]. Furthermore, in [13] it was shown that \mathcal{ML}_U is PSPACE-complete over general linear time.

Another open question is the complexity of $\mathcal{H}\mathcal{L}^{\downarrow, \oplus}$ -lin-SAT and $\mathcal{H}\mathcal{L}_{F,P}^{\downarrow}$ -lin-SAT (see the question marks in Table 1). Merely decidability follows from decidability of $\mathcal{H}\mathcal{L}_{F,P}^{\downarrow, \oplus}$ -lin-SAT [13].

Acknowledgments

We thank Balder ten Cate, Massimo Franceschet, Maarten Marx, and the anonymous referees for helpful suggestions, discussions, and comments.

References

1. Afanasiev, L., P. Blackburn, I. Dimitriou, B. Gaiffe, E. Goris, M. Marx and M. de Rijke, *PDL for ordered trees*, Journal of Applied Non-Classical Logics **15** (2005), pp. 115–135.
2. Areces, C., P. Blackburn and M. Marx, *A road-map on complexity for hybrid logics*, in: J. Flum and M. Rodríguez-Artalejo, editors, *CSL*, Lecture Notes in Computer Science **1683** (1999), pp. 307–321.
3. Areces, C., P. Blackburn and M. Marx, *The computational complexity of hybrid temporal logics*, Logic Journal of the IGPL **8** (2000), pp. 653–679.
4. Areces, C., P. Blackburn and M. Marx, *Hybrid logics: Characterization, interpolation and complexity*, J. Symb. Log. **66** (2001), pp. 977–1010.
5. Blackburn, P., *Representation, reasoning, and relational structures: a hybrid logic manifesto*, Logic Journal of the IGPL **8** (2000), pp. 339–365.
6. Blackburn, P., M. de Rijke and Y. Venema, “Modal Logic,” Cambridge Tracts in Theoretical Computer Science **53**, Cambridge University Press, 2001.
7. Blackburn, P. and J. Seligman, *Hybrid languages*, Journal of Logic, Language and Information **4** (1995), pp. 251–272.
8. Blackburn, P. and J. Seligman, *What are hybrid languages?*, in: M. Kracht, M. de Rijke, H. Wansing and M. Zakharyashev, editors, *Advances in Modal Logic*, CSLI Publications **1**, Stanford University, 1998 pp. 41–62.
9. Blackburn, P. and M. Tzakova, *Hybrid languages and temporal logic*, Logic Journal of the IGPL **7** (1999), pp. 27–54.
10. Börger, E., E. Grädel and Y. Gurevich, “The Classical Decision Problem,” Perspectives of Mathematical Logic, Springer-Verlag, 1997.
11. Dawar, A. and M. Otto, *Modal characterisation theorems over special classes of frames*, in: *LICS* (2005), pp. 21–30.
12. Franceschet, M. and M. de Rijke, *Model checking hybrid logics (with an application to semistructured data)*, Journal of Applied Logic (2005).
13. Franceschet, M., M. de Rijke and B.-H. Schlingloff, *Hybrid logics on linear structures: Expressivity and complexity*, in: *TIME* (2003), pp. 166–173.
14. Ganzinger, H., C. Meyer and M. Veanes, *The two-variable guarded fragment with transitive relations*, in: *Logic in Computer Science*, 1999, pp. 24–34.
15. Goranko, V., *Hierarchies of modal and temporal logics with reference pointers*, Journal of Logic, Language and Information **5** (1996), pp. 1–24.
16. Grädel, E., *On the restraining power of guards*, Journal of Symbolic Logic **64** (1999), pp. 1719–1742.
17. Grädel, E. and I. Walukiewicz, *Guarded fixed point logic*, in: *LICS*, 1999, pp. 45–54.

18. Immerman, N., A. M. Rabinovich, T. W. Reps, S. Sagiv and G. Yorsh, *The boundary between decidability and undecidability for transitive-closure logics*, in: J. Marcinkowski and A. Tarlecki, editors, *CSL*, Lecture Notes in Computer Science **3210** (2004), pp. 160–174.
19. Kieronski, E., *EXPSPACE-complete variant of guarded fragment with transitivity*, in: H. Alt and A. Ferreira, editors, *STACS*, Lecture Notes in Computer Science **2285** (2002), pp. 608–619.
20. Kieronski, E., *The two-variable guarded fragment with transitive guards is 2EXPTIME-hard*, in: A. D. Gordon, editor, *FoSSaCS*, Lecture Notes in Computer Science **2620** (2003), pp. 299–312.
21. Kracht, M., *Syntactic codes and grammar refinement*, Journal of Logic, Language and Information **4** (1995), pp. 41–60.
22. Kracht, M., *Inessential features*, in: C. Retoré, editor, *LACL*, Lecture Notes in Computer Science **1328** (1996), pp. 43–62.
23. Ladner, R. E., *The computational complexity of provability in systems of modal propositional logic*, SIAM J. Comput. **6** (1977), pp. 467–480.
24. Ono, H. and A. Nakamura, *On the size of refutation Kripke models for some linear modal and tense logics*, Studia Logica **34** (1980), pp. 325–333.
25. Papadimitriou, C. H., “Computational Complexity,” Addison-Wesley, 1994.
26. Prior, A., “Past, Present and Future,” Oxford University Press, 1967.
27. Reynolds, M., *The complexity of the temporal logic with “until” over general linear time*, J. Comput. Syst. Sci. **66** (2003), pp. 393–426.
28. Sistla, A. P. and E. M. Clarke, *The complexity of propositional linear temporal logics*, J. ACM **32** (1985), pp. 733–749.
29. Spaan, E., “Complexity of modal logics,” Ph.D. thesis, ILLC, University of Amsterdam (1993).
30. Spaan, E., *The complexity of propositional tense logics*, in: M. d. Rijke, editor, *Diamonds and Defaults*, Kluwer, 1993 pp. 287–307.
31. Szwast, W. and L. Tendera, *On the decision problem for the guarded fragment with transitivity*, in: *LICS*, 2001, pp. 147–156.
32. ten Cate, B. and M. Franceschet, *Guarded fragments with constants*, Journal of Logic, Language and Information **14** (2005), pp. 281–288.
33. ten Cate, B. and M. Franceschet, *On the complexity of hybrid logics with binders*, in: C.-H. L. Ong, editor, *CSL*, Lecture Notes in Computer Science **3634** (2005), pp. 339–354.

Appendix A

First-order Logic

First-order formulae are built from terms which are defined by $t ::= x \mid f(t_1, \dots, t_n)$, where x is a variable and f is an n -ary function symbol for $n \geq 0$. Formulae are defined by $\alpha ::= t_1 = t_2 \mid R(t_1, \dots, t_n) \mid \neg\alpha \mid \alpha \wedge \alpha' \mid \exists x \alpha$, where R is an n -ary relation symbol ($n \geq 1$), x is a variable, and t_i are terms. In the case of binary relations, we prefer the infix notation xRy . We use the standard abbreviations $\vee, \rightarrow, \leftrightarrow, \forall x \alpha, \top, \perp$ as well as the usual concepts of *free* and *bound* variables. To embed hybrid logic, it is sufficient to restrict the vocabulary to nullary function symbols (constants), unary and binary relation symbols only. We do not discuss further restrictions that are sufficient in this context.

FOL is interpreted in terms of *models* $\mathcal{M} = (D, I)$ consisting of a set D , the *domain*, and an *interpretation* I that maps constants to elements of D , *predicates* (unary relation symbols) to subsets of D , and binary relation symbols to binary relations over D . An *assignment* g is a function that maps each variable to an element of D .

The *satisfaction relation* ranges over models, assignments and formulae. If α contains exactly the free variables x_1, \dots, x_n , then α is denoted as $\alpha(x_1, \dots, x_n)$. The notation $\mathcal{M} \models \alpha[d_1, \dots, d_n]$, where $d_i \in D$, stands for “ \mathcal{M} satisfies α under the assignment that maps x_1, \dots, x_n to d_1, \dots, d_n , respectively.” This is defined as follows. $\mathcal{M} \models t_1 = t_2[d_1, \dots, d_n]$ iff the values of t_1 and t_2 are equal, where the value of a constant c is $I(c)$ and the value of a variable x_i is d_i . $\mathcal{M} \models P(t)[\dots]$ iff the value of t is in the set $I(P)$, analogously for binary relation symbols. The boolean cases are as usual, and finally, $\mathcal{M} \models \exists x \alpha[d_1, \dots, d_n]$ iff there exists some $d \in D$ such that $\mathcal{M} \models \alpha[d_1, \dots, d_n, x \mapsto d]$, where $[d_1, \dots, d_n, x \mapsto d]$ denotes the assignment obtained from $[d_1, \dots, d_n]$ by mapping x to d .

For *sentences* (i. e., formulae without free variables), satisfaction is independent of assignments. If a sentence α is satisfied in \mathcal{M} (under any assignment), we write $\mathcal{M} \models \alpha$ and say that α is *true* in \mathcal{M} . In general, a formula $\alpha(x_1, \dots, x_n)$ is *satisfiable* iff there exist a model \mathcal{M} and an appropriate assignment under which \mathcal{M} satisfies α .

Appendix B

Missing Parts of Section 3

B.1 Equivalence of \mathcal{HL}^\downarrow over Complete Frames and $\mathbf{MC}_=$

Before we present this connection precisely, we have to make a note on models. Every hybrid model can be viewed as a relational structure for its first-order correspondence language. This first-order language usually contains a binary relation to reflect the accessibility relation. For complete frames, the accessibility relation is trivial and can be ignored, respectively added when going from $\mathbf{MC}_=$ to \mathcal{HL}^\downarrow .

Lemma 16 *There are polynomial time functions mapping \mathcal{HL}^\downarrow formulas φ to $\mathbf{MC}_=$ formulas ψ and vice versa such that φ holds in a complete hybrid model \mathcal{M} if and only if ψ holds in the corresponding monadic structure.*

Proof. The mapping from \mathcal{HL}^\downarrow over complete frames to $\mathbf{MC}_=$ is based on the Standard Translation ST as defined in Section 2. The only rule of ST that uses the binary relation is the rule for the diamond operator: $\text{ST}_x(\Diamond\alpha) = \exists y(xRy \wedge \text{ST}_y(\alpha))$. But the right side can be reduced to $\exists y(\text{ST}_y(\alpha))$, since xRy always holds on complete frames.

Finally, we give rules for a reduction HT from $\mathbf{MC}_=$ to \mathcal{HL}^\downarrow .

$$\begin{aligned} \text{HT}(P(x)) &= \Diamond(x \wedge p), & \text{HT}(x = y) &= \Diamond(x \wedge y), \\ \text{HT}(\neg\varphi) &= \neg\text{HT}(\varphi), & \text{HT}(\varphi \wedge \psi) &= \text{HT}(\varphi) \wedge \text{HT}(\psi), \\ \text{HT}(\exists x.\varphi) &= \Diamond(\downarrow x.\text{HT}(\varphi)), \end{aligned}$$

Note that both mappings can be computed in polynomial time and do not blow up formula size. \square

B.2 Missing Proofs

Proof of Corollary 4. We can force a transitive frame, more precisely, the sub-frame generated by the current state, to be complete by adding $\downarrow x.\Box\Diamond x$. In this way, we can give a reduction from the satisfiability problem of $\mathbf{MC}_=$ by mapping a formula φ to $(\downarrow x.\Box\Diamond x) \wedge \text{HT}(\varphi)$. \square

Proof of Lemma 6. The proof is by induction on the structure of ψ . Most cases are trivial since s_1 is the only state that has to be considered, e.g., if $\psi = \psi_1 \wedge \psi_2$ we only need to know whether ψ_1 and ψ_2 hold at s_1 .

The interesting case is $\psi = \Diamond\xi$. Here, we need to know whether ξ holds at some state s' reachable from s_1 . This can only be affected by our substitution if s' is in the replaced subtree, and therefore strictly below s_1 . Thus, by our

assumption on g , all free variables in ψ are mapped to states different and not reachable from s' . We can conclude

$$\mathcal{B}, g, s' \models \xi \iff \mathcal{B}, g, s' \models \xi[free/\perp].$$

Hence ξ holds at some state in the replaced subtree rooted at u_{s_2} iff $\xi[free/\perp]$ is in the φ -type of u_{s_2} . The lemma follows because the new subtree has the same φ -type. \square

Proof of Lemma 7. Let φ be a $\mathcal{H}\mathcal{L}^\perp$ -sentence satisfiable over transitive frames and \mathcal{B}' a block tree model for φ .

An block tree satisfying the third condition can be obtained from \mathcal{B}' by applying Lemma 6. If there are two nodes u and v on a path, v below u , that have the same φ -type, we can replace the subtree rooted at u by the subtree rooted at v . This allows us to cut every finite repetition down to a single occurrence of a φ -type. The resulting structure is still a block tree.

Now, consider some node u and its φ -type t . For every sentence $\psi \in t$, we select some state in the subtree rooted at u such that ψ holds at this state. Let u_1, \dots, u_k be the nodes of those states, that are not in u . It easy to see, following similar tracks as in Lemma 6 that the block tree obtained by removing the nodes below u and inserting instead u_1, \dots, u_k as children of u is again a model of φ . Even more, the type of u does not change by this replacement.

By applying this argument top-down from the root, we get a block tree model for φ satisfying the first two conditions. The third one is not affected by this transformation. Note that some care is needed to make this approach work for infinite models. Basically, we must define a function that assigns to each node of the original model its set of witnesses. The resulting model is obtained by using this function in a straightforward fashion. \square

B.3 The Model Checking Algorithm MCFULL'

We have to modify the algorithm MCFULL of [12] in two points. First, we have to include the formulas resulting from replacing the free variables by \perp . The second point is the evaluation of formulas of the form $\Diamond\psi$, where we have to use the information about the states in C . The procedures are in the notation of [12], e.g., \mathbf{F} is used instead of \Diamond .

```

procedure CHECKF( $\mathcal{M}, g, \alpha$ )
  MCFULL( $\mathcal{M}, g, \alpha[free/\perp]$ )
  MCF( $\mathcal{M}, g, \alpha[free/\perp]$ )
  MCFULL( $\mathcal{M}, g, \alpha$ )
  MCF( $\mathcal{M}, g, \alpha$ )
end procedure

```

```

procedure  $\text{MC}_{\mathbf{F}}(\mathcal{M}, g, \alpha)$ 
  for  $w \in (L(\alpha) \cap M)$  do
    for  $v \in R^{-1}(w)$  do
       $L(\mathbf{F}\alpha, v) \leftarrow 1$ 
    end for
  end for
  for  $w \in L(\alpha[\text{free}/\perp])$  do
    for  $v \in R^{-1}(w)$  do
      if  $v \notin R(w)$  then
         $L(\mathbf{F}\alpha, v) \leftarrow 1$ 
      end if
    end for
  end for
end procedure

```

Appendix C

Missing Proof Details for Section 4

Proof of Lemma 9. It remains to prove the following claim: *For each formula α , α is satisfiable iff $f(\alpha)$ is satisfiable in some model that interprets R by a transitive relation.* Without loss of generality, we may assume that α has no free variables and that each variable is quantified exactly once. This can always be achieved by additional existential quantification and renaming, respectively.

“ \Rightarrow ”. Suppose α is satisfied by some model $\mathcal{M} = (D, I)$. From \mathcal{M} we construct a new model $\mathcal{M}^4 = (D^4, I^4)$, where $D^4 = D^0 \cup \dots \cup D^3$ and $D^i = \{d^i \mid d \in D\}$, for $i = 0, 1, 2, 3$. The interpretation I^4 is defined by $I^4(R) = \{(x^0, x^1), (x^2, x^1), (x^2, x^3), (y^0, x^3) \mid (x, y) \in I(R)\}$ and $I^4(P) = D^P$, $P = 0, 1, 2, 3$. $I^4(R)$ codes an $I(R)$ -transition from state x to y in \mathcal{M} as a sequence of backward and forward transitions from x_0 to y_0 via x_1, x_2, x_3 as shown in Figure 1. It is easy to see that $I^4(R)$ is transitive, since there is no state with incoming *and* outgoing $I^4(R)$ -edges.

We now show that for all subformulae $\beta(x_1, \dots, x_m)$ of α and all $d_1, \dots, d_m \in D$: $\mathcal{M} \models \beta[d_1, \dots, d_m]$ iff $\mathcal{M}^4 \models \beta^t[d_1^0, \dots, d_m^0]$. This immediately implies that \mathcal{M}^4 satisfies α^t .

We proceed by induction on β . The base case, $\beta = xRy$, is clear from the construction of $I^4(R)$. The boolean cases are obvious. For the case $\beta = \exists x \gamma$, we argue

$$\begin{aligned}
 \mathcal{M} \models \exists x \gamma[d_1, \dots, d_m] &\Leftrightarrow \exists d \in D (\mathcal{M} \models \gamma[d_1, \dots, d_m, x \mapsto d]) \\
 &\Leftrightarrow \exists d \in D (\mathcal{M}^4 \models \gamma^t[d_1^0, \dots, d_m^0, x \mapsto d^0]) \\
 &\Leftrightarrow \exists d' \in D^4 (\mathcal{M}^4 \models (0(x) \wedge \gamma^t)[d_1^0, \dots, d_m^0, x \mapsto d']) \\
 &\Leftrightarrow \mathcal{M}^4 \models \exists x (0(x) \wedge \gamma^t)[d_1^0, \dots, d_m^0] \\
 &\Leftrightarrow \mathcal{M}^4 \models (\exists x \gamma)^t[d_1^0, \dots, d_m^0].
 \end{aligned}$$

The first and second line are equivalent due to the induction hypothesis. The equivalence of the second and third line follows from the definition of R^4 ; for the direction from below to above, one must take into account that because x is interpreted by d' and $0(x)$ is satisfied, d' is indeed some d^0 .

“ \Leftarrow ”. Let $\mathcal{M} = (D, I)$ be a model satisfying α^t , where $I(R)$ is transitive. We construct a new model $\mathcal{M}' = (D', I')$, where $D' = I(0)$ and

$$\begin{aligned}
 I'(R) = \{ &(d, e) \in (D')^2 \mid \exists abc \in D ((d, a), (b, a), (b, c), (e, c) \in I(R) \\
 &\& a \in I(1) \& b \in I(2) \& c \in I(3)) \}.
 \end{aligned}$$

We now show that for all subformulae $\beta(x_1, \dots, x_m)$ of α and all $d_1, \dots, d_m \in D$: $\mathcal{M}' \models \beta[d_1, \dots, d_m]$ iff $\mathcal{M} \models \beta^t[d_1^0, \dots, d_m^0]$. This immediately implies that \mathcal{M}' satisfies α .

Again, the proof is via induction on β . The base case, $\beta = xRy$, is clear from the construction of $I'(R)$ and the fact that the translation of xRy requires $0(x)$ and $0(y)$. The boolean cases are obvious. For the case $\beta = \exists x \gamma$, we argue

$$\begin{aligned}
\mathcal{M}' \models \exists x \gamma [d_1, \dots, d_m] &\Leftrightarrow \exists d \in D' (\mathcal{M}' \models \gamma [d_1, \dots, d_m, x \mapsto d]) \\
&\Leftrightarrow \exists d \in I(0) (\mathcal{M} \models \gamma^t [d_1, \dots, d_m, x \mapsto d]) \\
&\Leftrightarrow \exists d \in D (\mathcal{M} \models (0(x) \wedge \gamma^t) [d_1, \dots, d_m, x \mapsto d]) \\
&\Leftrightarrow \mathcal{M} \models \exists x (0(x) \wedge \gamma^t) [d_1, \dots, d_m] \\
&\Leftrightarrow \mathcal{M} \models (\exists x \gamma)^t [d_1, \dots, d_m].
\end{aligned}$$

The induction hypothesis is applied between the first and second line. It is obvious that the second line implies the third; the backward direction is due to the fact that x is interpreted by d and $0(x)$ is satisfied, hence $d \in I(0)$.

This proves the above claim. \square

Proof of Theorem 10. In order to argue that each formula α is satisfiable iff $f(\alpha)$ is satisfiable, we assume w.l.o.g. that α is a sentence (see proof of Lemma 9). For the “ \Rightarrow ” direction, suppose α is satisfied by a model $\mathcal{M} = (D, I)$. By adding the spypoint s to D , we obtain the hybrid model $\mathcal{M}^h = (M^h, R^h, V^h)$, where $M^h = D \cup \{s\}$, $R^h = I(R) \cup \{(s, d) \mid d \in D\}$, and $V^h(p) = I(p)$. Clearly, \mathcal{M}^h satisfies $f(\alpha)$ at s —under any assignment, since $f(\alpha)$ is a sentence.

For the “ \Leftarrow ” direction, suppose $f(\alpha)$ is satisfied at state s of some hybrid model $\mathcal{M} = (M, R, V)$. The composition of $f(\alpha)$ enforces s to behave as the spypoint. It is easy to see that $\mathcal{M}' = (M - \{s\}, I)$, where $I(R) = R \upharpoonright_{M - \{s\}}$ and $I(P) = V(p)$, satisfies α .

In the case of $\mathcal{H}\mathcal{L}_{F,P}^\downarrow$, we must simulate the $@$ operator using P , which is possible in the presence of a spypoint and transitivity. We simply re-define $(\cdot)^t$ by $(xRy)^t = P(i \wedge F(x \wedge Fy))$, $(P(x))^t = P(i \wedge F(x \wedge p))$, $(\neg \alpha)^t = \neg(\alpha^t)$, $(\alpha \wedge \beta)^t = \alpha^t \wedge \beta^t$, and $(\exists x \alpha)^t = P(i \wedge F \downarrow x. \alpha^t)$. The rest of the proof is the same as for $\mathcal{H}\mathcal{L}_{F,P}^{\downarrow, @}$. \square

Proof of Theorem 11. Decidability immediately follows from decidability of $S\omega S$, using the Standard Translation ST. For the nonelementary lower bound, we reduce $\mathcal{H}\mathcal{L}_{F,P}^\downarrow$ -lin-SAT to $\mathcal{H}\mathcal{L}_{F,P}^\downarrow$ -tt-SAT and $\mathcal{H}\mathcal{L}^{\downarrow, @}$ -tt-SAT, respectively.

Let us first consider $\mathcal{H}\mathcal{L}_{F,P}^\downarrow$ -tt-SAT. A linear relation is a special case of a transitive tree, and we can in fact enforce that a transitive tree model is linear. All we have to do is to require that every point x has at most one direct successor, or, more elaborately, whenever x has some successor, then we name one of the *direct* successors y and ensure that *all* direct successors of x satisfy y . This translates as $\lambda = F\top \rightarrow F^1 \downarrow y. P^1 G^1 y$, where F^1 , P^1 , and G^1 refer to *direct* successors only. This can be expressed by means of U and S , for example $F^1 \varphi \equiv U(\varphi, \perp)$. Now, $U(\varphi, \psi)$ can be simulated by $\downarrow x. F(\varphi \wedge H(Px \rightarrow \psi))$; analogously for $S(\varphi, \psi)$.

Hence λ is expressible in our language and of constant length. An appropriate reduction function f is given by $f(\varphi) = \varphi \wedge \lambda \wedge \mathbf{H}\lambda \wedge \mathbf{H}\mathbf{G}\lambda$. It is easy to observe that φ is satisfiable in some linear model iff $f(\varphi)$ is satisfiable in some transitive tree.

In the case of $\mathcal{HL}^{\downarrow, @}$ -tt-SAT, we replace λ by $\lambda' = \Diamond \top \rightarrow \downarrow x. \Diamond^1 \downarrow y. @_x \Box^1 y$ and again express \Diamond^1 and \Box^1 by means of \mathbf{U} , which can be simulated as shown in Section 1. Instead of f , we use f' as a reduction function, where $f'(\varphi) = \varphi \wedge \lambda' \wedge \Box \lambda'$. \square

Appendix D

Missing Proof Details for Section 5

Proof of Theorem 12. It remains to prove the following two claims for each $\varphi \in \mathcal{ML}$.

- (1) If $\varphi \in \mathcal{ML}\text{-GLOBSAT}$, then $f(\varphi) \in \mathcal{ML}_{\text{U-tt-SAT}}$.
- (2) If $f(\varphi) \in \mathcal{ML}_{\text{U-trans-SAT}}$, then $\varphi \in \mathcal{ML}\text{-GLOBSAT}$.

Since each transitive tree is a transitive model, (1) and (2) imply the claim of this theorem.

(1). Suppose φ is satisfied in all states of some Kripke model $\mathcal{M} = (M, R, V)$. By considering the submodel generated by some arbitrary state, we can assume w.l.o.g. that \mathcal{M} has a root w_0 .

Due to the *tree model property* [6] there exists a tree-like model (a model whose underlying frame is a tree) that satisfies φ at all states, too. Hence we can suppose \mathcal{M} itself to be tree-like. From this model, we construct $\mathcal{M}' = (M, R^+, V)$, which is clearly a transitive tree.

Because of the tree-likeness of \mathcal{M} , we observe that for each pair $(w, v) \in R$, there exists no $u \in M$ between w and v in terms of R^+ , i.e. no u such that wR^+u and uR^+v . By means of this observation, we show that for all states $m \in M$ and all formulae $\psi \in \mathcal{ML}$: $\mathcal{M}, m \models \psi$ iff $\mathcal{M}', m \models \psi^t$. This claim implies that $\mathcal{M}', w_0 \models \varphi^t \wedge \Box \varphi^t$. It is proven by induction on the structure of ψ . The only interesting case is $\psi = \Diamond \vartheta$, and the necessary argument can be summarized as follows.

$$\begin{aligned}
 \mathcal{M}, m \models \Diamond \vartheta &\Leftrightarrow \exists n \in M (mRn \ \& \ \mathcal{M}, n \models \vartheta) \\
 &\Leftrightarrow \exists n \in M (mRn \ \& \ \mathcal{M}', n \models \vartheta^t) \\
 &\Leftrightarrow \exists n \in M (mR^+n \ \& \ \mathcal{M}', n \models \vartheta^t \ \& \ \neg \exists u \in M (mR^+u \ \& \ uR^+n)) \\
 &\Leftrightarrow \mathcal{M}', m \models \text{U}(\vartheta^t, \perp)
 \end{aligned}$$

In this argument, the equivalence of the first and the second line follows from the induction hypothesis. The second and third line are equivalent due to the above observation.

(2). Let $\mathcal{M} = (M, R, V)$ be a transitive model and $w_0 \in M$ such that $\mathcal{M}, w_0 \models f(\varphi)$. Again, we restrict ourselves to the submodel generated by w_0 . Hence all states of \mathcal{M} are accessible from w_0 .

Define a new Kripke model $\mathcal{M}' = (M, R', V)$ from \mathcal{M} , where $R' = \{(w, v) \in R \mid \neg \exists u \in M (wRuRv)\}$. We show that for all states $m \in M$ and all formulae $\psi \in \mathcal{ML}$: $\mathcal{M}', m \models \psi$ iff $\mathcal{M}, m \models \psi^t$. Again, we use induction on the structure of ψ with the only interesting case $\psi = \Diamond \vartheta$ and the following argument.

$$\begin{aligned}
 \mathcal{M}', m \models \Diamond \vartheta &\Leftrightarrow \exists n \in M (mR'n \ \& \ \mathcal{M}', n \models \vartheta)
 \end{aligned}$$

$$\begin{aligned} &\Leftrightarrow \exists n \in M ((n = w_0 \text{ or } w_0 R n) \ \& \ m R n \ \& \ \neg \exists u \in M (m R u R n) \ \& \ \mathcal{M}, n \models \vartheta^t) \\ &\Leftrightarrow \mathcal{M}, m \models U(\vartheta^t, \perp) \end{aligned}$$

The equivalence of the first and the second line is due to the fact that \mathcal{M} is rooted, the definition of R' as well as the induction hypothesis. Now, since $\mathcal{M}, w_0 \models \varphi^t \wedge \Box \varphi^t$, we conclude that for all states $x \in M$, $\mathcal{M}, x \models \varphi^t$. The previous claim implies that \mathcal{M}' satisfies φ at all states. \square

Proof of Lemma 13. Either problem can be reduced to the other via a simple bijection $f : \mathcal{HL}_{U,S}^{\textcircled{U}} \rightarrow \mathcal{HL}_{U^{++},S^{++}}^{\textcircled{U}}$ or its inverse, respectively. This function simply replaces every occurrence of U (or S , respectively) in the input formula by U^{++} (or S^{++} , respectively). Obviously, f and f^{-1} can be computed in polynomial time. It is straightforward to inductively verify the following two propositions.

- (1) For every $\varphi \in \mathcal{HL}_{U,S}^{\textcircled{U}}$: If φ is satisfied in a state m of some transitive model \mathcal{M} , then $\mathcal{M}, m \models f(\varphi)$.
- (2) For all $\varphi \in \mathcal{HL}_{U^{++},S^{++}}^{\textcircled{U}}$: If φ is satisfied in a state m of some model $\mathcal{M} = (M, R, V)$, then the transitive model $\mathcal{M}' = (M, R^+V)$ satisfies $f^{-1}(\varphi)$ at m . \square

Proof of Theorem 14. We first embed $\mathcal{HL}_{U^{++},S^{++}}^{\textcircled{U}}$ into the loosely μ -guarded fragment μLGF of first-order logic [17]. Since the satisfiability problem for μLGF -sentences is 2EXP-complete [17], we obtain a 2EXP upper bound for $\mathcal{HL}_{U,S}^{\textcircled{U}}$ -trans-SAT by Lemma 13. As a second step, we will show a reduction from $\mathcal{HL}_{U,S}^E$ -trans-SAT to $\mathcal{HL}_{U,S}^{\textcircled{U}}$ -trans-SAT.

For the embedding into μLGF , we enhance the Standard Translation ST (see Section 2) by the rule

$$\text{ST}_x(U^{++}(\varphi, \psi)) = \exists y [xR^+y \wedge \text{ST}_y(\varphi) \wedge \forall z ((xR^+z \wedge zR^+y) \rightarrow \text{ST}_z(\psi))]$$

and an analogous rule for $\text{ST}_x(S^{++}(\varphi, \psi))$. ST_y and ST_z are defined by exchanging x, y, z cyclically.

It remains to take care of the R^+ expressions. But xR^+y can be expressed by $[\text{LFP } W(x, y). (xRy \vee \exists z (zRy \wedge xWz))]xy$, yielding a μLGF -sentence with three variables. (If U^{++} operators are nested, variables can be “recycled”.) The constants from the translations of nominals can be eliminated introducing new variables as shown in [16]. The whole translation only requires time polynomial in the length of the input formula.

As for the more expressive language with E , we can embed the stronger language $\mathcal{HL}_{U,S}^E$ into $\mathcal{HL}_{U,S}^{\textcircled{U}}$ using a spypoint argument (see [7, 2]) and exploiting the fact that we are restricted to transitive frames. By adding a spypoint (i.e. a point s that sees all other points and is named by the fresh nominal i) to a transitive model, $E\varphi$ can be simulated by $@_i \Diamond \varphi$. Hence, if we take the translation $(\cdot)^t : \mathcal{HL}_{U,S}^E \rightarrow \mathcal{HL}_{U,S}^{\textcircled{U}}$ that simply replaces all occurrences of E as shown, we

obtain a reduction function $f : \mathcal{H}\mathcal{L}_{U,S}^E\text{-trans-SAT} \rightarrow \mathcal{H}\mathcal{L}_{U,S}^{\textcircled{E}}\text{-trans-SAT}$ by setting $f(\varphi) = i \wedge \neg \Diamond i \wedge \Diamond \varphi^t$.

Clearly, f is computable in polynomial time. It is straightforward to verify that f is an appropriate reduction function: If $\varphi \in \mathcal{H}\mathcal{L}_{U,S}^E$ is satisfied at some point of some transitive model, add the spypoint s and the according accessibilities. The new model satisfies $f(\varphi)$ at s . For the converse, if a transitive model satisfies $f(\varphi)$ at some point s , then s must be a spypoint, and φ^t is satisfied at another point m . Remove s and the according accessibilities from the model and observe that it now satisfies φ at m . \square

A note on the discrepancy between the upper and lower bound for $\mathcal{H}\mathcal{L}_{U,S}^E\text{-trans-SAT}$. Since the 2EXP result for μLGF in [17] holds for sentences without constants only, constants — which arise from the translation of nominals — must be reformulated using new variables. This causes an unbounded number of variables in the first-order vocabulary, because we have no restriction on the number of nominals in our hybrid language.

Could we assume that the number of nominals were bounded, then the described reduction would yield guarded fixpoint sentences of bounded width. In this case, satisfiability is EXP-complete [17]. It is not known whether in the case of a bounded number of variables, but an arbitrary number of constants, satisfiability for μLGF -sentences also decreases from 2EXP to EXP, as it is the case for the fragment without the μ operator [32]. If there were a positive answer to this question, an EXP upper bound for our satisfiability problem would follow.

Proof of Theorem 15. Since we want to give an embedding into $\mathcal{PDL}_{\text{tree}}$, we first introduce syntax and semantics of $\mathcal{PDL}_{\text{tree}}$.

$\mathcal{PDL}_{\text{tree}}$ is the language of propositional dynamic logic with four atomic programs **left**, **right**, **up**, and **down** that are associated with the relations “left sister”, “right sister”, “parent”, and “daughter” in trees. It consists of all formulae of the form $\varphi ::= p \mid \neg\varphi \mid \varphi \wedge \varphi' \mid \langle \pi \rangle \varphi$, where $p \in \text{ATOM}$ and π is a program. Programs are defined by $\pi ::= \text{left} \mid \text{right} \mid \text{up} \mid \text{down} \mid \pi; \pi' \mid \pi \cup \pi' \mid \pi^* \mid \varphi?$, where φ is a formula. We abbreviate $[\pi]\varphi := \neg\langle \pi \rangle \neg\varphi$ and $a^+ := a; a^*$ for atomic programs a .

A $\mathcal{PDL}_{\text{tree}}$ model is a multi-modal model $\mathcal{M} = (T, R_{\text{down}}, R_{\text{right}}, V)$, where T is a finite tree with an order relation on all immediate successors of any node, R_{down} is the successor relation and R_{right} is the “next-sister” relation. The set of relations is extended to arbitrary programs as follows: $R_{\text{up}} = R_{\text{down}}^-$, $R_{\text{left}} = R_{\text{right}}^-$, $R_{\pi; \pi'} = R_{\pi} \circ R_{\pi'}$, $R_{\pi \cup \pi'} = R_{\pi} \cup R_{\pi'}$, $R_{\pi^*} = R_{\pi}^*$, and $R_{\varphi?} = \{(m, m) \mid \mathcal{M}, m \models \varphi\}$. The satisfaction relation for atomic formulae and booleans is defined as for hybrid logic. The modal case is given by $\mathcal{M}, m \models \langle \pi \rangle \varphi$ iff $\exists n \in T(m R_{\pi} n \ \& \ \mathcal{M}, n \models \varphi)$. A formula φ is *satisfiable* if there exists a model $\mathcal{M} = (T, R_{\text{down}}, R_{\text{right}}, V)$, such that $\mathcal{M}, m \models \varphi$, where m is the root of T . For any φ , let $N(\varphi)$ be the set of all nominals occurring in φ .

We reduce $\mathcal{H}\mathcal{L}_{U,S}^E\text{-trans-SAT}$ to $\mathcal{PDL}_{\text{tree}}\text{-SAT}$ and define a translation $(\cdot)^t : \mathcal{H}\mathcal{L}_{U,S}^E \rightarrow \mathcal{PDL}_{\text{tree}}$ by $p^t = p$, $p \in \text{ATOM}$ (nominals are translated into atomic

propositions); $(\neg\varphi)^t = \neg(\varphi^t)$; $(\varphi \wedge \psi)^t = \varphi^t \wedge \psi^t$; $(E\varphi)^t = \langle \mathbf{up}^*; \mathbf{down}^* \rangle \varphi^t$; $(U(\varphi, \psi))^t = \langle (\mathbf{down}; \psi^t?)*; \mathbf{down} \rangle \varphi^t$; and $(S(\varphi, \psi))^t = \langle (\mathbf{up}; \psi^t?)*; \mathbf{up} \rangle \varphi^t$.

Since $\mathcal{PDL}_{\text{tree}}$ has no nominals, we must enforce that (the translation of) each nominal is true at exactly one point by requiring

$$\begin{aligned} \nu(i) = \langle \mathbf{down}^* \rangle i \wedge [\mathbf{down}^*] \Big(i \rightarrow ([\mathbf{down}^+] \neg i \wedge [\mathbf{up}^+] \neg i \\ \wedge [\mathbf{up}^*; \mathbf{left}^+; \mathbf{down}^*] \neg i \wedge [\mathbf{up}^*; \mathbf{right}^+; \mathbf{down}^*] \neg i) \Big) \end{aligned}$$

to hold for each nominal i . As a reduction function, we have $f(\varphi) = \langle \mathbf{down}^* \rangle \varphi^t \wedge \bigwedge_{i \in N(\varphi)} \nu(i)$.

It is clear that f is computable in polynomial time and straightforward to show that f is an appropriate reduction function: Suppose, φ is satisfiable in some finite transitive tree model $\mathcal{M} = (M, R, V)$ based on the tree (M, R') with root w . Then $f(\varphi)$ is satisfiable in w of the $\mathcal{PDL}_{\text{tree}}$ model based on the tree (M, R') , equipped with the valuation V . For the converse, if $f(\varphi)$ is satisfied at the root of some $\mathcal{PDL}_{\text{tree}}$ model $\mathcal{M} = (M, R_{\text{down}}, R_{\text{right}}, V)$, then φ^t is true at some point w , and each nominal is true at exactly one point of \mathcal{M} . Hence $(M, R_{\text{down}}^+, V)$ —where R_{down}^+ is the transitive closure of R_{down} —is a hybrid transitive tree model satisfying φ at w .

Now there is one drawback in the reduction via f . According to our definition of a tree, it is not necessary that a (transitive) tree is finite or has a root. A node can have infinitely many successors, or there may be an infinitely long forward or backward path from some point. For most practical applications these cases are certainly hardly of interest, but we strive for a more general result. If we do allow for infinite depth or width, the above translation into $\mathcal{PDL}_{\text{tree}}$ —which is interpreted over finite, rooted trees—is not sufficient.

To overcome finiteness, it suffices to re-examine the proof for the EXP upper bound of $\mathcal{PDL}_{\text{tree}}$ -satisfiability in [1]. This proof in fact covers a more general result, too, namely that satisfiability of $\mathcal{PDL}_{\text{tree}}$ formulae over (not necessarily finite) trees is in EXP.

To cater for the fact that “our” trees do not need to have roots, we first observe that satisfiability over *rooted* transitive trees is reducible to satisfiability over (arbitrary) transitive trees, because a root is expressible by $\mathbf{PH}\perp$ in our language. Since the lower bound from Theorem 12 holds with respect to rooted transitive trees, it also holds for arbitrary ones.

In order to obtain the upper bound with respect to arbitrary transitive trees, we propose a modification of the above reduction via f . The basic idea is to turn the backward path from the node w (that is to satisfy φ) into a forward path, such that w becomes the root of the transformed model. Thus all predecessors of w (and *their* predecessors) become successors and must be marked by a fresh proposition \flat . (See Figure 2.)




Fig. 2. Making predecessors successors.

As a first step, we construct a new translation $(\cdot)^{tb}$ from $(\cdot)^t$ retaining all but the U/S-cases. For U/S, we replace all occurrences of the programs **down** and **up** by programs that incorporate the new structure and the fact that for b -nodes, their predecessors used to be their successors, and their b -successors used to be their predecessors. We define

$$\begin{aligned} (U(\varphi, \psi))^{tb} &= \langle (\mathbf{dn}'; \psi^*)^*; \mathbf{dn}' \rangle \varphi, \quad \text{where } \mathbf{dn}' = (\mathbf{down}; \neg b?) \cup (b?; \mathbf{up}), \quad \text{and} \\ (S(\varphi, \psi))^{tb} &= \langle (\mathbf{up}'; \psi^*)^*; \mathbf{up}' \rangle \varphi, \quad \text{where } \mathbf{up}' = (\neg b?; \mathbf{up}) \cup (b?; \mathbf{down}; b?). \end{aligned}$$

Note that we do not change the translation of $E\varphi$. The only thing that remains to do is to enforce that there is exactly one path at whose every node b is true. This means that b must be true at the root node and at exactly one successor of each node satisfying b . This can be expressed by

$$\begin{aligned} \beta &= b \wedge [\mathbf{down}^*] \left(b \rightarrow ([\mathbf{left}^+] \neg b \wedge [\mathbf{right}^+] \neg b \wedge \langle \mathbf{down} \rangle b) \right) \\ &\quad \wedge [\mathbf{down}^*] (\neg b \rightarrow [\mathbf{down}] \neg b). \end{aligned}$$

It is now straightforward to show that f^b , given by $f^b(\varphi) = \varphi^{tb} \wedge \beta \wedge \bigwedge_{i \in N(\varphi)} \nu(i)$, is indeed an appropriate reduction function.

(Note that φ^{tb} replaces $\langle \mathbf{down}^* \rangle \varphi^{tb}$, because we have turned w into the new root node.) \square